

JAMES SCOTT

scottjab@gmail.com GitHub @scottjab (937) 239-4054

STAFF INFRASTRUCTURE ENGINEER, SITE RELIABILITY ENGINEER

Profile

A curious Staff Software Engineer with 16 years of experience and a infrastructure background. Skilled at scaling out systems from 0 users to millions. Self driven and believes that though iteration most difficult problems can be solved.

Skills

Reliable Observable
Operations: Systems that
can be seen can be fixed.

Distributed System
Design: Designed
hardened production
systems that handle
mission critical work loads
at global scales.

Self Directed: Has worked
remote for 7 years, and
has a proven record to
deliver results.

Technical

Go

Python

Java

Scala

Ruby

Linux

Kubernetes

Sequins

Puppet

Chef

Kafka

mongodb

hbase

cassandra

git

hadoop

hive

Dataproc

Experience

Squarespace (Remote)

Tech Lead

June 2019-Aug 2022

Started as the Tech Lead for the MongoDB subteam on the Databases team. The first main task was scaling the large sharded MongoDB clusters. Migrated backups to a snapshot system in GCP, and worked with engineering teams to fix data models that did not scale.

I ended my tenure at Squarespace as the Tech Lead for Data Infrastructure. My team was responsible for all streaming and data processing infrastructure. In a short time we were able to migrate two legacy on-premises aging Hadoop clusters to Dataproc in GCP. During this migration I worked with many teams to burn down technical debt and migrate many ETLs and reports to offerings based in GCP.

Stripe (Remote)

Software Engineer

Apr 2016-Feb 2019

Worked with the Storage Team to globally scale out MongoDB, by writing tools to observe how the organization used the database, including code to fingerprint queries to optimize index usage. Also wrote tooling to automate day to day operations on a large set of sharded mongo clusters. In addition to the MongoDB work, I was in charge of development of Stripe internal open source data warehouse named "Sequins". This distributed read only key/value store sat at the end of the ETL pipeline at Stripe and stored everything from counts to trained machine learning models.

Urban Airship (Portland, OR)

Lead Operations Engineer (Security)

Mar 2013-Mar 2016

Scaled Urban Airship's push notification and analytic platforms through the mobile boom of the early 2010's. Leveraged Kafka to successfully deliver millions of messages a day on real hardware. I built deployment and operations tooling using python, and JVM based languages such as clojure and java. I was also leading the company's efforts around observability. I designed and built a distributed scalable log parsing and analysis system on top of kafka, that was API compatible with the Elasticsearch/Logstash/Kibana system.

University of Dayton (Dayton, OH)

Systems Engineer

Jul 2006-Mar 2013

I was responsible for the University's successful single sign on Identity management project. This included a single sign on for all university resources. I also worked on the in house DNS/DHCP self registration system. Over the years the university transitioned from Netware to Linux to serve file and print services, during this transition I lead the project to move from hand configured nodes to using centralized configuration management with puppet. This greatly reduced toil for the operators.
